

Imutils

Tools for blazingly fast statistical analysis

Josef Graf

McMaster University

Table of contents

1. RARity Case Study
2. Welcome to Imutils
 - Getting Started
 - Working with Files
 - Configuration
 - Matrix Manipulation
 - Statistical Functions
 - Analysis Functions
3. plots.r
4. Conclusion

RARity Case Study

- The original code was written in **148 lines** of pure R, and by no fault of the original author (purely that of R) was slow.
- A single phenotype and gene block took **42 minutes and 38 seconds** to run.
- For a full UKB run this would take days, if not weeks.

The Code

```
lmutils::set_num_worker_threads(32)
phenos <- lmutils::Mat$new("phenos.rkyv")
phenos$remove_column("eid")
genos <- lmutils::Mat$new("genos.rkyv")
genos$min_column_sum(2)
genos$na_to_column_mean()
genos$standardize_columns()
genos$remove_column("eid")
df <- lmutils::calculate_r2(genos, phenos)
```

- The new code is only **9 lines** long and makes use of a variety of functions from `lmutils`.
- A single phenotype and gene block can now be run in as little as **44 seconds**.

What!?! But how!

- Parallelism
 - Using 1 thread instead of 32 only adds about 4 and a half minutes.
- rkyv
 - Loading from `.RData` instead of `.rkyv` only adds about 10 seconds.
- So where does the rest of the time go?
- Rust
- faer
 - SIMD, parallelism, linear algebra magic.

Welcome to Imutils

Installation

First, you need to install the Rust toolchain. You can do this by running

```
curl --proto '=https' --tlsv1.2 -sSf https://sh.rustup.rs | sh
```

Then, you can install lmutils by running the following in R. Once the package is in a more final state I'm going to submit it to CRAN.

```
install.packages(  
  "https://github.com/GMELab/lmutils.r/archive/refs/heads/master.tar.gz",  
  repos=NULL) # use .zip for Windows  
# OR  
devtools::install_github("GMELab/lmutils.r")
```


Introductory definitions

- **Matrix-convertible object:** a data frame, matrix, file name, numeric column vector, or a Mat object.
- **List of matrix-convertible objects:** a list of matrix-convertible objects, a character vector of file names, or a single matrix-convertible object.
- **Standard output file:** a character vector or list of file names matching the length of the inputs or NULL to return the output. If a single input (not in a list) was given, the output will not be in a list.
- **Core parallelism:** the number of primary operations that will be run in parallel (i.e. gene blocks processed at once, files read at once, etc.).

The Mat object

- The Mat object is a wrapper around lmutils's internal representation of a matrix.
- It allows for easy, efficient, and lazy manipulation of the matrix.
- It is the recommended way to interact with the newest version of lmutils.

```
mat <- lmutils::Mat$new("phenos.rkyv")  
eids <- mat$col("eid")  
r <- mat$r() # just like any other R matrix  
colnames(r) == mat$colnames() # TRUE  
r$eid == eids # TRUE
```

lmutils supports a variety of both compressed and uncompressed file formats.

- `.csv` (`.csv.gz`) comma-separated values, first row is the column header.
- `.tsv` (`.tsv.gz`) tab-separated values, first row is the column header.
- `.txt` (`.txt.gz`) space-separated values, first row is the column header.
- `.rkyv` (`.rkyv.gz`) a binary format that is both fast and space-efficient (**recommended**).
- `.RData` an R-specific binary format.

Reading and Writing Files

lmutils supports two methods of reading and writing files, using the Mat object or global functions.

```
mat <- lmutils::Mat$new("phenos.rkyv")  
r <- mat$r()  
mat$save("phenos.csv")
```

both functions support a list of files

```
lmutils::load("phenos.rkyv")  
lmutils::save(  
  list("file1.csv", matrix(1:9, nrow=3), 1:3,  
    data.frame(a=1:3, b=4:6), mat),  
  c("file1.tsv", "file2.rkyv.gz", "file3.rkyv",  
    "file4.rdata", "file5.tx.gz")),  
)
```

Configuration

lmutils provides a number of global configuration options that can be set using environment variables or through various functions. Please note, these generally cannot be changed after calling any other lmutils functions.

- `LMUTILS_LOG` / `lmutils::set_log_level`: the log level, defaults to `info`. The valid levels (in order of increasing verbosity) are `off`, `error`, `warn`, `info`, `debug`, and `trace`.
- `LMUTILS_CORE_PARALLELISM` / `lmutils::set_core_parallelism`: the number of primary operations that will be run in parallel, defaults to 16.
- `LMUTILS_NUM_WORKER_THREADS` / `lmutils::set_num_worker_threads`: the number of worker threads, defaults to `num_cpus` / 2.
- `LMUTILS_ENABLE_PREDICTED` / `lmutils::disable_predicted` / `lmutils::enable_predicted`: whether to calculate and return predicted values in linear models, defaults to disabled for performance reasons.

Matrix Manipulation

Imutils provides more than 30 functions just for manipulating Mat objects, not to mention more than a dozen other global functions for matrices and data frames. Here are a few examples.

- `Mat$remove_column(col)`: remove a column by name.
- `Mat$min_column_sum(threshold)`: remove columns with a sum below a threshold.
- `Mat$na_to_column_mean()`: replace NA values with the column mean.
- `Mat$standardize_columns()`: standardize columns to have a mean of 0 and a standard deviation of 1.
- `Mat$transpose()`: transpose the matrix.
- `Mat$sort(col_idx)`: sort the matrix by a column.
- `Mat$dedup(col_idx)`: deduplicate the matrix by a column.

A full list of functions can be found in the documentation on GitHub at github.com/GMELab/Imutils.r.

lmutils also provides a number of lightning fast statistical functions.

- `lmutils::compute_r2(v1, v2)`: compute the R^2 between two vectors.
- `lmutils::mean(v)`: compute the mean of a vector.
- `lmutils::median(v)`: compute the median of a vector.
- `lmutils::sd(v)`: compute the standard deviation of a vector.
- `lmutils::var(v)`: compute the variance of a vector.

calculate_r2

The first kind of primary analysis is calculating the R^2 value between a list of matrix convertible objects (i.e. gene blocks) and a single matrix convertible object (i.e. phenotypes).

```
df <- lmutils::calculate_r2(list("genos1.rkyv", "genos2.rkyv",  
"genos3.rkyv"), "phenos.rkyv")
```

The output is a data frame with the columns:

- r2: the R^2 value.
- adj_r2: the adjusted R^2 value.
- data: the name or index of the data.
- outcome: the name or index of the outcome.
- n: the number of samples, i.e. rows.
- m: the number of predictors, i.e. SNPs.
- predicted: the predicted values of the outcome for each sample, only returned if enabled.

The second kind of primary analysis is calculating the p -values of the each pair of data and outcome columns. In some testing it can do a 20 thousand column block in about 2-3 seconds.

```
df <- lmutils::column_p_values(list("genos1.rkyv",  
"genos2.rkyv", "genos3.rkyv"), "phenos.rkyv")
```

The output is a data frame with the columns:

- `p_value`: the p -value of the regression.
- `beta`: the slope of the regression.
- `intercept`: the intercept of the regression.
- `data`: the name or index of the data.
- `data_column`: the index of the column in the data.
- `outcome`: the name or index of the outcome.

plots.r

- `plots.r` is a small package for creating publication-ready plots.
- It currently only supports forest plots thanks to `metafor`, but I'm happy to add more upon request!

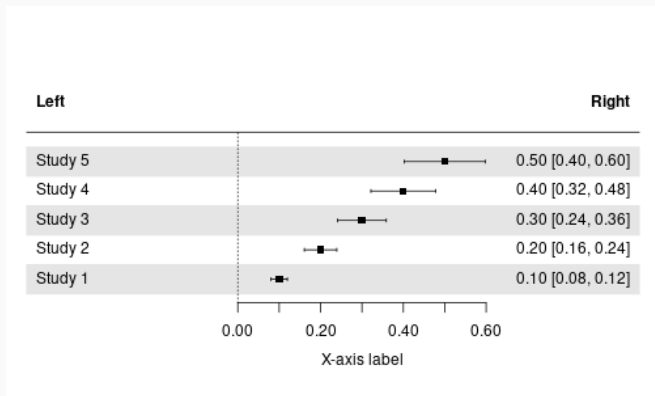
It can be installed by running the following in R.

```
install.packages(  
  "https://github.com/GMELab/plots.r/archive/refs/heads/master.tar.gz",  
  repos=NULL) # use .zip for Windows  
# OR  
devtools::install_github("GMELab/plots.r")
```

Basic Forest Plot

```
plots.r::basic_forest_plot(  
  x = c(0.1, 0.2, 0.3, 0.4, 0.5),  
  se = c(0.01, 0.02, 0.03, 0.04, 0.05),  
  width = 500,  
  height = 300,  
  name = "Example",  
  header = c("Left", "Right"),  
  slab = c("Study 1", "Study 2", "Study 3",  
           "Study 4", "Study 5"),  
  xlab = "X-axis label",  
)
```

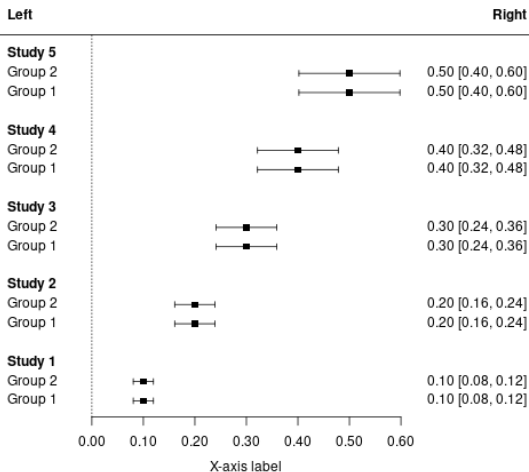
Basic Forest Plot



Grouped Forest Plot

```
plots.r::grouped_forest_plot(  
  x = list(c(0.1, 0.2, 0.3, 0.4, 0.5),  
           c(0.1, 0.2, 0.3, 0.4, 0.5)),  
  se = list(c(0.01, 0.02, 0.03, 0.04, 0.05),  
            c(0.01, 0.02, 0.03, 0.04, 0.05)),  
  width = 500,  
  height = 500,  
  name = "Grouped_Example",  
  header = c("Left", "Right"),  
  slab = c("Study 1", "Study 2", "Study 3",  
           "Study 4", "Study 5"),  
  xlab = "X-axis label",  
  glab = c("Group 1", "Group 2"),  
)
```

Grouped Forest Plot



Conclusion

Conclusion

- Thank you for listening!
- Feel free to ask me questions anytime, I'm always happy to help or add new things. :)
- grafj1@mcmaster.ca or Josef Graf on Teams.